

# Lecture 5, September 29, 2025

## Forecasting using ARIMA and ADL models

Marta Boczon

Department of Economics

Copenhagen Business School

mbo.eco@cbs.dk

```
In [1]: #install.packages("quantmod")
#install.packages("fredr")
#install.packages("ggfortify")
#install.packages('urca')
#install.packages("tseries")
#install.packages("forecast")
#install.packages("dynlm")
#install.packages("stargazer")
#install.packages("pracma")
#install.packages("dLagM")
#install.packages("gets")
#install.packages("car")
#install.packages("lmtest")
#install.packages("astsa")
#install.packages("zoo")
#install.packages("conflicted")
#library(conflicted)
#conflict_prefer("lag", "stats")
options(warn=-1)
```

### Introduction

Today we will focus on finding **ADL models** and forecasting using **ARIMA** and **ADL** models. For this exercise, we will work with two Canadian macroeconomic series:

1. **Canadian GDP** – the same quarterly series we introduced in Lecture 3.
2. **Weekly hours worked in the Canadian manufacturing sector** – the same quarterly series we introduced in Lecture 4.

### Data Transformations

We begin by importing the data and transforming it to achieve **stationarity**, which is a prerequisite for applying ARIMA and/or ADL models.

- For **GDP**, we take the **first difference of the logarithm**.
  - This transformation allows us to interpret the series in terms of *quarterly growth rates*, while also stabilizing the variance and removing the stochastic trend.
- For **hours worked**, we take the **first difference of the level series**.
  - This transformation removes the unit root identified in the original series.

With these transformations, we now have two stationary series suitable for ARIMA and ADL modeling.

```
In [2]: # Import the dataset from a CSV file downloaded from FRED
gdp = read.csv("NGDPRSAXDCCAQ.csv")
# Convert the second column of the dataset into a time series object
tsgdp = ts(gdp[, 2], frequency = 4, start = c(1961, 1))
# Transform the series into (approximate) percentage changes
dgdp = diff(log(tsgdp))
```

```
In [3]: # Import the dataset from a CSV file downloaded from FRED
hours = read.csv("HOHWMN02CAM065N.csv")
# Convert the second column of the dataset into a time series object
tshours = ts(hours[, 2], frequency = 4, start = c(1960, 1))
# Transform the series into a change from the previous period
dhours = diff(tshours)
```

### Exhaustive ADL Subset Search Based on Parameter Significance

```
In [41]: #####
## Exhaustive ADL Subset Search based on parameter significance
## Goal: start from a general ADL(4,4) and iteratively drop the
##       Least significant regressor until all remaining have
```

```

##  $|t| \geq 1.96$  ( $\approx 5\%$  two-sided level), or none remain.
#####
library(ts)
library(gets)

## ---- 1) Create lagged variables -----
## Assume 'data' is a bivariate ts object:
## - column 1: y (e.g., GDP growth)
## - column 2: x (e.g., ΔHours worked)
## We build up to 4 lags for each series. Using negative 'k' in lag(z, -k)
## shifts the series BACK by k steps so that, e.g., Lag(y, -1) corresponds to  $y_{t-1}$ .
y0 = data[, 1]           #  $y_t$  (dependent variable at time t)

y1 = lag(data[, 1], -1)   #  $y_{t-1}$  (first own lag of y)
y2 = lag(data[, 1], -2)   #  $y_{t-2}$ 
y3 = lag(data[, 1], -3)   #  $y_{t-3}$ 
y4 = lag(data[, 1], -4)   #  $y_{t-4}$ 

x1 = lag(data[, 2], -1)   #  $x_{t-1}$  (first lag of x)
x2 = lag(data[, 2], -2)   #  $x_{t-2}$ 
x3 = lag(data[, 2], -3)   #  $x_{t-3}$ 
x4 = lag(data[, 2], -4)   #  $x_{t-4}$ 

tempdata = ts.intersect(    # Align  $y_t$  and all lags so rows line up (drop initial NA rows from lags)
  y0,
  y1, y2, y3, y4,
  x1, x2, x3, x4
)

# Separate dependent and independent variables after alignment
depvar = tempdata[, 1]      #  $y_t$ 
indvar = tempdata[, 2:ncol(tempdata)]  # [ $y_{t-1..4}$ ,  $x_{t-1..4}$ ]

# Name columns explicitly for readability/tracking during elimination
colnames(indvar) = c("y1", "y2", "y3", "y4", "x1", "x2", "x3", "x4")

# Initialize min | $t|$  to trigger entry into the loop
tmin = 0

# Main backward-elimination loop based on individual t-statistics:
# Continue while the smallest absolute t-stat among non-constant coefficients is < 1.96.
while (tmin < 1.96) {

  # Estimate ARX with current regressor set 'indvar' and a constant (mc=TRUE).
  # try(..., silent=TRUE) avoids stopping on errors (e.g., singular design), and
  # suppressWarnings(...) silences routine warnings during re-estimation.
  model = try(suppressWarnings(arx(depvar, mxreg = indvar, mc = TRUE))), silent = TRUE)

  print(model)          # Print iteration's model object (useful to monitor retained vars)

  b = coef(model)       # Named vector of estimated coefficients (incl. constant 'mconst')
  V = vcov(model)       # Variance-covariance matrix of estimates
  se = sqrt(diag(V))   # Standard errors from diagonal of V

  tstat = b / se        # Compute coefficient-wise t-statistics

  # Drop the constant's t-stat from the elimination rule:
  # We only compare *slope* coefficients (lags) to the threshold, not 'mconst'.
  tstat = tstat[setdiff(names(tstat), "mconst")]

  if (length(tstat) == 0) break      # If no slope coefficients are left, stop the loop

  tstat = abs(tstat)               # Work with absolute t-values for two-sided testing
  tmin = min(tstat)                # Record the smallest | $t|$  among current regressors

  # Identify the variable with the minimum | $t|$  and prepare a new regressor set that excludes it.
  # setdiff(colnames(indvar), names(which.min(tstat))) returns all columns except the one with min | $t|$ .
  k = setdiff(colnames(indvar), names(which.min(tstat)))

  # If dropping that variable leaves no columns, set indvar to NULL
  # (intercept-only ARX on next iteration/exit).
  if (length(k) == 0) indvar = NULL

  # Otherwise, keep all remaining columns
  # (drop = FALSE preserves matrix structure when selecting a single column).
  if (length(k) != 0) indvar = indvar[, k, drop = FALSE]

  # Loop continues until every remaining slope has | $t| \geq 1.96$  (or no slopes remain).
}

print(model)

```

Date: Mon Sep 29 14:46:22 2025  
Dependent var.: depvar  
Method: Ordinary Least Squares (OLS)  
Variance-Covariance: Ordinary  
No. of observations (mean eq.): 246  
Sample: 1962(2) to 2023(3)

Mean equation:

	coef	std.error	t-stat	p-value
mconst	0.00663890	0.00131206	5.0599	8.411e-07 ***
y1	-0.01557032	0.06819366	-0.2283	0.81959
y2	0.03205575	0.06793092	0.4719	0.63744
y3	0.01339174	0.06757077	0.1982	0.84307
y4	0.06839600	0.06566642	1.0416	0.29867
x1	0.00552215	0.00254866	2.1667	0.03126 *
x2	-0.00225605	0.00270668	-0.8335	0.40540
x3	0.00263759	0.00267562	0.9858	0.32524
x4	0.00073847	0.00255386	0.2892	0.77271
---				

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	1.2137e-05	1	0.9972
Ljung-Box ARCH(1)	4.5282e+01	1	1.706e-11 ***
---			

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

SE of regression 0.01252  
R-squared 0.04534  
Log-lik.(n=246) 732.93535

Date: Mon Sep 29 14:46:22 2025  
Dependent var.: depvar  
Method: Ordinary Least Squares (OLS)  
Variance-Covariance: Ordinary  
No. of observations (mean eq.): 246  
Sample: 1962(2) to 2023(3)

Mean equation:

	coef	std.error	t-stat	p-value
mconst	0.0067450	0.0011953	5.6427	4.737e-08 ***
y1	-0.0149218	0.0679775	-0.2195	0.82644
y2	0.0316242	0.0677588	0.4667	0.64113
y4	0.0680705	0.0655133	1.0390	0.29984
x1	0.0055132	0.0025431	2.1679	0.03116 *
x2	-0.0022083	0.0026905	-0.8208	0.41260
x3	0.0027989	0.0025436	1.1004	0.27228
x4	0.0008553	0.0024799	0.3449	0.73048
---				

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	1.0435e-05	1	0.9974
Ljung-Box ARCH(1)	4.4983e+01	1	1.988e-11 ***
---			

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

SE of regression 0.01250  
R-squared 0.04519  
Log-lik.(n=246) 732.93286

Date: Mon Sep 29 14:46:22 2025  
Dependent var.: depvar  
Method: Ordinary Least Squares (OLS)  
Variance-Covariance: Ordinary  
No. of observations (mean eq.): 246  
Sample: 1962(2) to 2023(3)

Mean equation:

	coef	std.error	t-stat	p-value
mconst	0.00662607	0.00106327	6.2318	2.066e-09 ***
y2	0.03231252	0.06755135	0.4783	0.63285
y4	0.06770278	0.06536129	1.0358	0.30133
x1	0.00533163	0.00239997	2.2215	0.02725 *
x2	-0.00235540	0.00260046	-0.9058	0.36597
x3	0.00276942	0.00253502	1.0925	0.27573
x4	0.00077061	0.00244479	0.3152	0.75288
---				

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	0.046023	1	0.8301
Ljung-Box ARCH(1)	46.959257	1	7.248e-12 ***
---			

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1  
SE of regression 0.01247  
R-squared 0.04499  
Log-lik.(n=246) 732.92369

Date: Mon Sep 29 14:46:22 2025  
Dependent var.: depvar  
Method: Ordinary Least Squares (OLS)  
Variance-Covariance: Ordinary  
No. of observations (mean eq.): 246  
Sample: 1962(2) to 2023(3)

Mean equation:

	coef	std.error	t-stat	p-value
mconst	0.0065696	0.0010461	6.2801	1.574e-09 ***
y2	0.0322440	0.0674241	0.4782	0.63292
y4	0.0725581	0.0634010	1.1444	0.25359
x1	0.0051510	0.0023262	2.2144	0.02774 *
x2	-0.0025801	0.0024961	-1.0337	0.30233
x3	0.0025007	0.0023829	1.0494	0.29503
---				

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	0.041735	1	0.8381
Ljung-Box ARCH(1)	46.579415	1	8.798e-12 ***
---			

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

SE of regression 0.01245  
R-squared 0.04460  
Log-lik.(n=246) 732.88614

Date: Mon Sep 29 14:46:22 2025  
Dependent var.: depvar  
Method: Ordinary Least Squares (OLS)  
Variance-Covariance: Ordinary  
No. of observations (mean eq.): 246  
Sample: 1962(2) to 2023(3)

Mean equation:

	coef	std.error	t-stat	p-value
mconst	0.00680178	0.00092512	7.3523	3.026e-12 ***
y4	0.07524938	0.06304960	1.1935	0.23385
x1	0.00528289	0.00230611	2.2908	0.02284 *
x2	-0.00218588	0.00235226	-0.9293	0.35368
x3	0.00278639	0.00230311	1.2098	0.22753
---				

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	0.050922	1	0.8215
Ljung-Box ARCH(1)	45.269191	1	1.717e-11 ***
---			

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

SE of regression 0.01243  
R-squared 0.04369  
Log-lik.(n=246) 732.78042

Date: Mon Sep 29 14:46:22 2025  
Dependent var.: depvar  
Method: Ordinary Least Squares (OLS)  
Variance-Covariance: Ordinary  
No. of observations (mean eq.): 246  
Sample: 1962(2) to 2023(3)

Mean equation:

	coef	std.error	t-stat	p-value
mconst	0.00678932	0.00092477	7.3417	3.198e-12 ***
y4	0.08350118	0.06240354	1.3381	0.182125
x1	0.00591930	0.00220145	2.6888	0.007669 **
x3	0.00343547	0.00219402	1.5658	0.118695
---				

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	0.10175	1	0.7497
Ljung-Box ARCH(1)	47.45418	1	5.63e-12 ***
---			

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

SE of regression 0.01243

```
R-squared      0.04026
Log-lik.(n=246) 732.34979
```

```
Date: Mon Sep 29 14:46:22 2025
Dependent var.: depvar
Method: Ordinary Least Squares (OLS)
Variance-Covariance: Ordinary
No. of observations (mean eq.): 246
Sample: 1962(2) to 2023(3)
```

Mean equation:

```
    coef std.error t-stat p-value
mconst 0.00742438 0.00079498 9.3391 < 2.2e-16 ***
x1      0.00588612 0.00220488 2.6696  0.008107 **
x3      0.00339291 0.00219736 1.5441  0.123868
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Diagnostics and fit:

```
    Chi-sq df  p-value
Ljung-Box AR(1)  0.04785  1  0.8268
Ljung-Box ARCH(1) 47.38444  1 5.834e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
SE of regression  0.01245
R-squared         0.03316
Log-lik.(n=246)  731.45033
```

```
Date: Mon Sep 29 14:46:22 2025
Dependent var.: depvar
Method: Ordinary Least Squares (OLS)
Variance-Covariance: Ordinary
No. of observations (mean eq.): 246
Sample: 1962(2) to 2023(3)
```

Mean equation:

```
    coef std.error t-stat p-value
mconst 0.0073650 0.0007963 9.2490 < 2e-16 ***
x1      0.0052970 0.0021778 2.4323  0.01572 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Diagnostics and fit:

```
    Chi-sq df  p-value
Ljung-Box AR(1)  0.043719  1  0.8344
Ljung-Box ARCH(1) 49.375064  1 2.114e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
SE of regression  0.01248
R-squared         0.02367
Log-lik.(n=246)  730.25453
```

```
Date: Mon Sep 29 14:46:22 2025
Dependent var.: depvar
Method: Ordinary Least Squares (OLS)
Variance-Covariance: Ordinary
No. of observations (mean eq.): 246
Sample: 1962(2) to 2023(3)
```

Mean equation:

```
    coef std.error t-stat p-value
mconst 0.0073650 0.0007963 9.2490 < 2e-16 ***
x1      0.0052970 0.0021778 2.4323  0.01572 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Diagnostics and fit:

```
    Chi-sq df  p-value
Ljung-Box AR(1)  0.043719  1  0.8344
Ljung-Box ARCH(1) 49.375064  1 2.114e-12 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
SE of regression  0.01248
R-squared         0.02367
Log-lik.(n=246)  730.25453
```

```
In [5]: library(gets)
#https://search.r-project.org/CRAN/refmans/sysid/html/arx.html
#https://cran.r-project.org/web/packages/getsm/getsm.pdf

##GETS-modelling of mean specification:

#getsm(object, t.pval=0.05, wald.pval=t.pval, vcov.type=NULL,
#do.pet=TRUE, ar.LjungB=list(lag=NULL, pval=0.025),
```

```

#arch.LjungB=NULL, pval=0.025), normality.JarqueB=NULL,
#user.diagnostics=NULL, info.method=c("sc","aic","aicc", "hq"),
#gof.function=NULL, gof.method=NULL, keep=NULL, include.gum=FALSE,
#include.1cut=TRUE, include.empty=FALSE, max.paths=NULL, tol=1e-07,
#turbo=FALSE, print.searchinfo=TRUE, plot=NULL, alarm=FALSE)

# y-series: current value and its lags ----

y0 = data[,1]                      # y_t : dependent variable at time t (first column of 'data')

y1 = lag(data[,1], -1)              # y_{t-1} : first lag of y
y2 = lag(data[,1], -2)              # y_{t-2} : second lag of y
y3 = lag(data[,1], -3)              # y_{t-3} : third lag of y
y4 = lag(data[,1], -4)              # y_{t-4} : fourth lag of y

# x-series: lags ----

x1 = lag(data[,2], -1)              # x_{t-1} : first lag of x (second column of 'data')
x2 = lag(data[,2], -2)              # x_{t-2} : second lag of x
x3 = lag(data[,2], -3)              # x_{t-3} : third lag of x
x4 = lag(data[,2], -4)              # x_{t-4} : fourth lag of x

# Align all series on a common time index (drop rows with any NA from Lagging) ----

tempdata = ts.intersect(
  y0,                                # y_t
  y1, y2, y3, y4,                     # y Lags: y_{t-1..4}
  x1, x2, x3, x4                     # x Lags: x_{t-1..4}
)

# Split into dependent and regressor blocks ----

depvar = tempdata[, 1]                # dependent variable y_t after alignment
indvar = tempdata[, 2:ncol(tempdata)]  # matrix of regressors: (y_{t-1..4}, x_{t-1..4})

# Estimate general ADL with a constant, then run gets reduction ----

model.temp = arx(depvar, mxreg = indvar, mc=TRUE) # ARX/ADL: depvar ~ indvar + constant

getfit = getsm(
  model.temp,
  arch.Ljung = NULL,                 # (diagnostics option) here set to NULL → skip ARCH Ljung-Box check
  print.searchinfo = FALSE,           # suppress verbose search output
  t.pval = 0.05                      # 5% threshold for t-tests during general-to-specific reduction
)

getfit                               # final selected (simplified) model object

```

Date: Mon Sep 29 13:30:11 2025  
 Dependent var.: depvar  
 Method: Ordinary Least Squares (OLS)  
 Variance-Covariance: Ordinary  
 No. of observations (mean eq.): 246  
 Sample: 1962(2) to 2023(3)

SPECIFIC mean equation:

```

  coef std.error t-stat p-value
mconst 0.0073650 0.0007963 9.2490 < 2e-16 ***
x1      0.0052970 0.0021778 2.4323 0.01572 *
---
Signif. codes:  0 '****' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	0.043719	1	0.8344
Ljung-Box ARCH(1)	49.375064	1	2.114e-12 ***
SE of regression	0.01248		
R-squared	0.02367		
Log-lik.(n=246)	730.25453		

The manual loop performs a **backward-elimination search** starting from a general **ADL(4,4)** model. In each iteration, it estimates the current model with a constant, computes  $t$ -statistics for all non-constant coefficients, and identifies the regressor with the smallest absolute  $t$ -statistic. That weakest term is dropped, the model is re-estimated on the reduced set, and the process repeats.

The loop stops when either no regressors remain or all surviving slope coefficients satisfy  $|t| \geq 1.96$ , which corresponds to an approximate **5% two-sided significance threshold** under large-sample normality.

This greedy, nested-model approach is transparent and easy to follow, though it does not guarantee a **globally optimal subset** in terms of information criteria.

The `getsm()` routine from the **GETS** framework automates the same reduction while allowing for optional diagnostic checks.

## Exercise 1

Use **iterative elimination of insignificant parameters** to find the best model for **Canadian GDP** as a function of **weekly hours worked in the manufacturing sector in Canada** and **personal consumption expenditures (PCE) in the US**. This approach begins from a general specification and gradually removes regressors that do not meet a chosen significance threshold, leaving a **parsimonious model** where all remaining coefficients are individually meaningful.

To transform the monthly data on PCE into quarterly frequency, as we practiced in **Lecture 3**, please run the code below. This step ensures that the frequency of all variables is aligned before estimating the model, which is crucial for meaningful inference.

```
In [6]: library(zoo)

# Import the dataset from a CSV file downloaded from FRED -----
temp = read.csv("PCE.csv")           # reads the CSV into a data.frame 'temp'

# Parse monthly dates -----
# Convert the character/date column 'observation_date' to a monthly index.
# as.yearmon stores dates as Year-Month (e.g., 2020-01) without a specific day.
temp$date = as.yearmon(temp$observation_date)

# Aggregate monthly to quarterly -----
# Collapse the monthly series 'PCE' to calendar quarters using the mean of the
# three months in each quarter. Replace FUN as needed (sum, Last, etc.).
pce = aggregate(PCE ~ as.yearqtr(Date), data = temp, FUN = mean)

head(pce)                         # quick peek at the first few quarterly observations
tail(pce)                          # quick peek at the most recent quarterly observations

# Convert the second column of the quarterly data to a ts object -----
# Create a quarterly 'ts' starting in 1959 Q1. 'frequency = 4' indicates quarterly data.
# pce[, 2] is the numeric series (PCE Levels) and the first column is the quarter index.
tspce = ts(pce[, 2], frequency = 4, start = c(1959, 1))
```

A data.frame: 6 × 2

	as.yearqtr(Date)	PCE
	<yearqtr>	<dbl>
1	1959 Q1	309.4667
2	1959 Q2	315.5000
3	1959 Q3	320.7333
4	1959 Q4	322.8667
5	1960 Q1	326.3667
6	1960 Q2	332.2000

A data.frame: 6 × 2

	as.yearqtr(Date)	PCE
	<yearqtr>	<dbl>
262	2024 Q2	19682.70
263	2024 Q3	19938.47
264	2024 Q4	20255.47
265	2025 Q1	20461.60
266	2025 Q2	20643.90
267	2025 Q3	20802.00

## Solution to Exercise 1

```
In [7]: # Calculate the number of observations in the time series object
nd = length(tspce)

# Print a message along with the number of observations
cat("There are", nd, "observations in our data set.")
```

There are 267 observations in our data set.

```
In [8]: # Load required libraries
library(ggplot2)  # for plotting
library(ggfortify) # for autoplot of time series objects

# Set default figure size (width = 8, height = 8)
options(repr.plot.width = 8, repr.plot.height = 8)
```

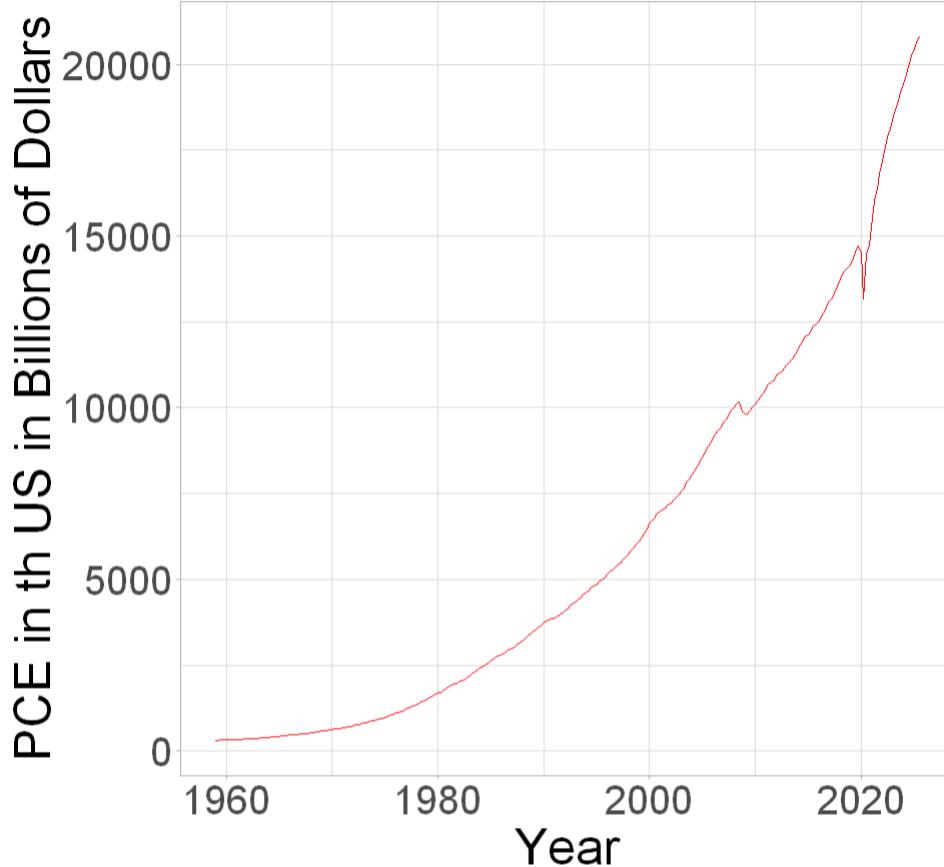
```

# Create an initial time series plot of tspce in red
fig = autoplot(tspce, colour = 'red')

# Customize the plot
fig = fig +
  theme(aspect.ratio = 1) + # make plot square-shaped
  theme_light() + # use a light theme
  theme(aspect.ratio = 1) + # reinforce square aspect ratio
  theme(plot.margin = ggplot2::margin(0.2, 0.2, 0.2, 0.2, "cm")) + # adjust margins
  theme(text = element_text(size = 30)) + # set large font size for readability
  labs(x = "Year") + # Label x-axis
  labs(y = "PCE in th US in Billions of Dollars")# Label y-axis

# Display the plot
fig

```



In [9]: `library(urca)  
summary(ur.df(tspce, type='trend', lags=8, selectlags="AIC"))`

```

#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####

Test regression trend

Call:
lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)

Residuals:
    Min      1Q      Median      3Q      Max 
-1527.58   -12.04    -0.07    16.55  1023.51 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 12.868353  25.266869  0.509   0.6110    
z.lag.1     0.013244  0.005267  2.514   0.0125 *  
tt        -0.058207  0.375801 -0.155   0.8770    
z.diff.lag -0.080146  0.063596 -1.260   0.2087    
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 136.6 on 254 degrees of freedom
Multiple R-squared:  0.18,    Adjusted R-squared:  0.1703 
F-statistic: 18.58 on 3 and 254 DF,  p-value: 6.27e-11

```

Value of test-statistic is: 2.5144 34.6555 25.5984

```

Critical values for test statistics:
      1pct  5pct 10pct
tau3 -3.98 -3.42 -3.13
phi2  6.15  4.71  4.05
phi3  8.34  6.30  5.36

```

## Interpretation

We compare the reported test statistics (**2.5144**, **34.6555**, **25.5984**) to the critical values table.

Since the test statistic exceeds the relevant critical values, we **reject the null hypothesis** that the deterministic trend and constant are both zero.

This means that even after including a deterministic trend and constant, the series still behaves as if it has a **unit root** (i.e., it is **non-stationary**).

Because the data show **exponential growth** (a trending level effect), we need to first take the **logarithm** of the series. This stabilizes the trend by converting exponential growth into a *linear trend*. After logging, we then **difference the series** (log-differences correspond to growth rates) and retest for a unit root on the transformed series.

```
In [10]: # Transform the PCE time series into growth rates -----
# 1. log(tspce): take natural logarithm of the level series (stabilizes exponential growth)
# 2. diff(...): compute first differences of the Logged series (log-differences = growth rates)
dpce = diff(log(tspce))
```

```
In [11]: library(urca)
summary(ur.df(dpce, type='trend', lags=8, selectlags="AIC"))

#####
# Augmented Dickey-Fuller Test Unit Root Test #
#####

Test regression trend
```

```
Call:
lm(formula = z.diff ~ z.lag.1 + 1 + tt + z.diff.lag)

Residuals:
    Min      1Q  Median      3Q     Max 
-0.106581 -0.003592 -0.000344  0.003573  0.086624 

Coefficients:
            Estimate Std. Error t value Pr(>|t|)    
(Intercept) 1.679e-02 2.693e-03  6.236 1.88e-09 *** 
z.lag.1     -7.276e-01 9.985e-02 -7.287 4.07e-12 *** 
tt          -3.752e-05 1.063e-05 -3.531 0.000492 *** 
z.diff.lag1 -2.774e-01 8.568e-02 -3.238 0.001365 **  
z.diff.lag2 -1.410e-01 6.171e-02 -2.284 0.023206 *   
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.01131 on 252 degrees of freedom
Multiple R-squared:  0.5113,   Adjusted R-squared:  0.5036 
F-statistic: 65.92 on 4 and 252 DF,  p-value: < 2.2e-16
```

Value of test-statistic is: -7.2868 17.7311 26.5967

```
Critical values for test statistics:
  1pct  5pct 10pct
tau3 -3.98 -3.42 -3.13
phi2  6.15  4.71  4.05
phi3  8.34  6.30  5.36
```

## Interpretation

The reported test statistic for the unit root test is **-7.2868**, which is far below the **1% critical value** of **-3.98**. This means we can clearly **reject the null hypothesis** of a unit root in the differenced (growth rate) series. In other words, there is **no second unit root**.

However, the presence of **significant deterministic components** is still indicated. This suggests that while the growth rate series is stationary in terms of stochastic trends, it still contains a **deterministic trend**.

Therefore, we need to **detrend the growth rates** to remove this deterministic trend and obtain a properly **stationary series** for further modeling.

```
In [34]: # Load the pracma package (provides the detrend() function)
library(pracma)

# Remove the Linear trend from the series "temp"
# as.numeric(temp) ensures pracma::detrend() works
# (expects a numeric vector, not a ts object)
# 'Linear' specifies linear detrending
temp = detrend(as.numeric(dpce), 'linear')

# Define the start date for the time series
dpce = ts(temp, frequency = 4, start = c(1959, 2))
```

```
In [13]: library(gets)
#https://search.r-project.org/CRAN/refmans/sysid/html/arx.html
#https://cran.r-project.org/web/packages/getts/getts.pdf
```

```

##GETS-modelling of mean specification:

#getsrm(object, t.pval=0.05, wald.pval=t.pval, vcov.type=NULL,
#do.pet=TRUE, ar.LjungB=list(lag=NULL, pval=0.025),
#arch.LjungB=list(lag=NULL, pval=0.025), normality.JarqueB=NULL,
#user.diagnostics=NULL, info.method=c("sc", "aic", "aicc", "hq"),
#gof.function=NULL, gof.method=NULL, keep=NULL, include.gum=FALSE,
#include.1cut=TRUE, include.empty=FALSE, max.paths=NULL, tol=1e-07,
#turbo=FALSE, print.searchinfo=TRUE, plot=NULL, alarm=FALSE)

# y-series: current value and its lags ----

data = ts.intersect(dgdp, dhours, dpce)

y0 = data[,1]           # y_t : dependent variable at time t
# (first column of 'data')

y1 = lag(data[,1], -1)   # y_{t-1} : first lag of y
y2 = lag(data[,1], -2)   # y_{t-2} : second lag of y
y3 = lag(data[,1], -3)   # y_{t-3} : third lag of y
y4 = lag(data[,1], -4)   # y_{t-4} : fourth lag of y

# x-series: lags ----

x1 = lag(data[,2], -1)   # x_{t-1} : first lag of x (second column of 'data')
x2 = lag(data[,2], -2)   # x_{t-2} : second lag of x
x3 = lag(data[,2], -3)   # x_{t-3} : third lag of x
x4 = lag(data[,2], -4)   # x_{t-4} : fourth lag of x

# z-series: lags ----

z1 = lag(data[,3], -1)   # z_{t-1} : first lag of z (third column of 'data')
z2 = lag(data[,3], -2)   # z_{t-2} : second lag of z
z3 = lag(data[,3], -3)   # z_{t-3} : third lag of z
z4 = lag(data[,3], -4)   # z_{t-4} : fourth lag of z

# Align all series on a common time index (drop rows with any NA from Lagging) ----

tempdata = ts.intersect(
  y0,                      # y_t
  y1, y2, y3, y4,          # y Lags: y_{t-1..4}
  x1, x2, x3, x4,          # x Lags: x_{t-1..4}
  z1, z2, z3, z4           # z Lags: z_{t-1..4}
)

# Split into dependent and regressor blocks ----

depvar = tempdata[, 1]      # dependent variable y_t after alignment
indvar = tempdata[, 2:ncol(tempdata)] # matrix of regressors:
# (y_{t-1..4}, x_{t-1..4}, z_{t-1..4})

# Estimate general ADL with a constant, then run gets reduction ----

model.temp = arx(depvar, mxreg = indvar, mc=TRUE) # ARX/ADL: depvar ~ indvar + constant

getfit = getsm(
  model.temp,
  arch.Ljung = NULL,        # (diagnostics option) here set to NULL → skip ARCH Ljung-Box check
  print.searchinfo = FALSE, # suppress verbose search output
  t.pval = 0.05             # 5% threshold for t-tests during general-to-specific reduction
)

getfit                         # final selected (simplified) model object

```

Date: Mon Sep 29 13:30:12 2025  
 Dependent var.: depvar  
 Method: Ordinary Least Squares (OLS)  
 Variance-Covariance: Ordinary  
 No. of observations (mean eq.): 246  
 Sample: 1962(2) to 2023(3)

SPECIFIC mean equation:

```

  coef std.error t-stat p-value
mconst 0.0073650 0.0007963 9.2490 < 2e-16 ***
x1      0.0052970 0.0021778 2.4323 0.01572 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

```

Diagnostics and fit:

	Chi-sq	df	p-value
Ljung-Box AR(1)	0.043719	1	0.8344
Ljung-Box ARCH(1)	49.375064	1	2.114e-12 ***
---			
Signif. codes:	0 '***'	0.001 '**'	0.01 '*'
	0.05 '.'	0.1 ' '	1

```

SE of regression  0.01248
R-squared         0.02367
Log-lik.(n=246)  730.25453

```

## Conclusion

### Does US PCE help to model Canadian GDP?

From the **GETS** output for the specific mean equation, the final model retains only the **intercept** and one **non-PCE regressor (x1)**, while all lags of **US PCE** were eliminated by the reduction search because their *t*-statistics were below the **5% significance threshold**.

In other words, once we control for **weekly hours** (and the constant), the **US PCE terms** do not add incremental explanatory power for **Canadian GDP** — so they are not included in the **parsimonious specification**.

## Forecasting ARIMA Models

### AR(1) Model

**Note:** We will forecast **AR(1)** and **ARIMA(1,0,1)** models *for illustration*, so we can match R's built-in functions with our manual computations. As we saw last week, the **best model by BIC** was **white noise**, while the **best model by AIC** was a **subset ARIMA** with an AR term at lag 3 and MA terms at lags 3 and 4.

**R: Forecasting GDP growth rate using AR(1), 1-step ahead:**

$$y_t = \alpha_0 + \alpha_1 y_{t-1} + \epsilon_t$$

```
In [35]: library(stargazer)
# Fit an AR(1) to already-stationary GDP growth (dgdp)
model.a = arima(dgdp, order = c(1, 0, 0))

# Print the model
stargazer(model.a, type = "text")
```

```
=====
Dependent variable:
-----
dgdp
-----
ar1          0.027
(0.063)

intercept    0.007*** 
(0.001)

-----
Observations      256
Log Likelihood   758.550
sigma2          0.0002
Akaike Inf. Crit. -1,511.100
-----
Note:           *p<0.1; **p<0.05; ***p<0.01
```

```
In [36]: # Load forecasting
library(forecast)

# One-step-ahead forecast from the fitted ARIMA model
# h = 1 → forecast the next period only
# Level sets the prediction interval(s); add/remove as you like
f.a = forecast::forecast(model.a, h = 1, level = c(80, 95))

# Print a tidy summary (point forecast + intervals)
f.a

# Access pieces programmatically
f.a$mean      # point forecast
f.a$lower     # Lower bounds for Levels specified
f.a$upper     # upper bounds for Levels specified
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2025 Q2	0.007406762	-0.008612644	0.02342617	-0.01709281	0.03190633

A Time Series: 1 × 1

**Qtr2**

**2025** 0.007406762

A Time Series: 1 × 2

	80%	95%
--	-----	-----

**2025 Q2** -0.008612644 -0.01709281

A Time Series: 1 × 2

	80%	95%
--	-----	-----

**2025 Q2** 0.02342617 0.03190633

```
In [16]: # --- Manual calculations for 1-step-ahead AR(1) forecast from model.a ---
# Extract estimated coefficients from the fitted ARIMA model
```

```

coef.a = coef(model.a)

# Reparameterize: for an AR(1) with mean  $\mu$ , intercept ( $c$ ) =  $\mu*(1 - \varphi)$ 
alpha0 = coef.a["intercept"] * (1 - coef.a["ar1"])
alpha1 = coef.a["ar1"]

# Extract last observed value of the stationary series  $y_T$ 
yT = tail(dgdp, n = 1)

# Compute the conditional mean for  $T+1$ 
yT1.a = alpha0 + alpha1 * yT

# Express as ~percent (dgdp is  $\Delta \log \Rightarrow$  approx growth rate), rounded to two decimal places
yT1.a = round(yT1.a * 100, 2)

# 1-step-ahead point forecast:
cat("AR(1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is", as.numeric(yT1.a), "%.")

```

AR(1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is 0.74 %.

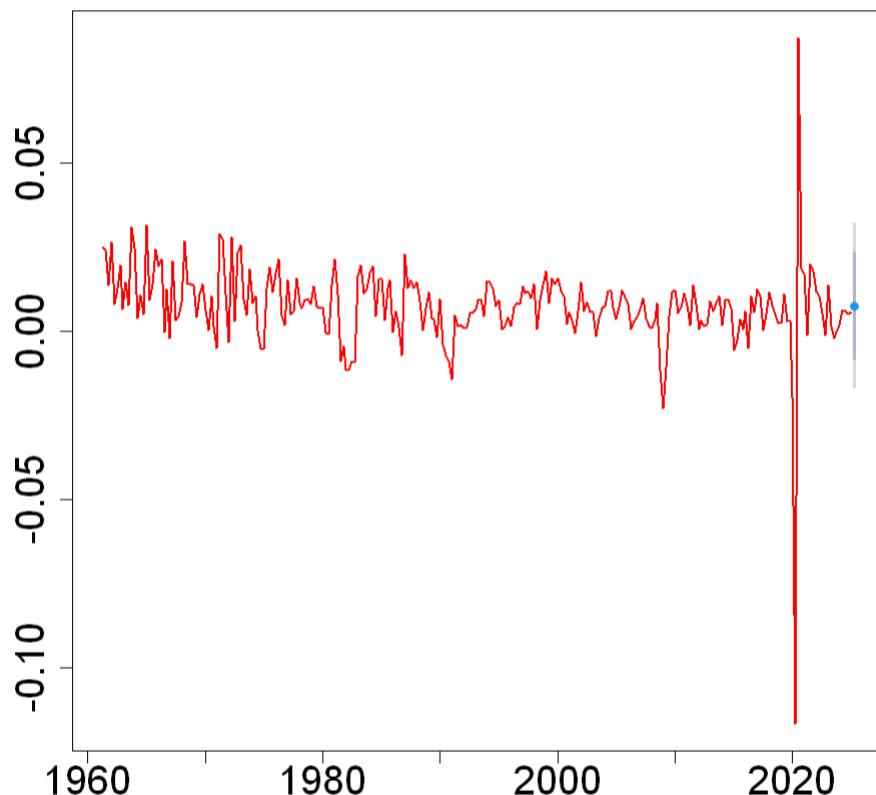
```

In [17]: # Set Jupyter/IRkernel figure size in inches for the output cell
options(repr.plot.width = 8, repr.plot.height = 8)

# Enlarge base R graphics text: axis labels, tick labels, and (potential) main title
par(cex.lab = 2, cex.axis = 2, cex.main = 2)

# Plot the 1-step-ahead forecast object from the 'forecast' package.
plot(f.a, aspect.ratio = 1, col = "red", lwd = 2, main = "")

```



## Forecasting ARIMA Models

### ARIMA(1,1) Model

R: Forecasting GDP growth rate using ARIMA(1,1), 1-step ahead:

$$y_t = \alpha_0 + \alpha_1 y_{t-1} + \delta_1 \epsilon_{t-1} + \epsilon_t$$

```

In [18]: library(stargazer)
# Fit an ARIMA(1,1) to already-stationary GDP growth (dgdp)
model.b = arima(dgdp, order = c(1, 0, 1))

# Print the model
stargazer(model.b, type = "text")

```

```
=====
Dependent variable:
-----
          dgdp
-----
ar1           0.990***  

              (0.015)  

  
ma1          -0.964***  

              (0.024)  

  
intercept     0.008***  

              (0.002)  

  
-----
Observations      256  

Log Likelihood   760.936  

sigma2            0.0002  

Akaike Inf. Crit. -1,513.873
=====
```

Note: \*p<0.1; \*\*p<0.05; \*\*\*p<0.01

In [19]:

```
# Load forecasting
library(forecast)

# One-step-ahead forecast from the fitted ARIMA model
# h = 1 → forecast the next period only
# level sets the prediction interval(s); add/remove as you like
f.b = forecast::forecast(model.b, h = 1, level = c(80, 95))

# Print the output
f.b

# Access pieces programmatically
f.b$mean      # point forecast
f.b$lower     # Lower bounds for levels specified
f.b$upper     # upper bounds for levels specified
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2025 Q2	0.005713515	-0.01014489	0.02157192	-0.01853982	0.02996685

A Time Series: 1 × 1

**Qtr2**

**2025** 0.005713515

A Time Series: 1 × 2

80%	95%
-----	-----

**2025 Q2** -0.01014489 -0.01853982

A Time Series: 1 × 2

80%	95%
-----	-----

**2025 Q2** 0.02157192 0.02996685

In [20]:

```
# Manual calculations
coef.b = coef(model.b)

alpha0 = coef.b["intercept"] * (1 - coef.b["ar1"])
alpha1 = coef.b["ar1"]
delta1 = coef.b["ma1"]

# Extract last observed value of the stationary series y_T and the corresponding residual e_T
yT = tail(dgdp, n = 1)
eT = tail(residuals(model.b), n = 1)

# Compute the conditional mean for T+1
yT1.b = alpha0 + alpha1 * yT + delta1 * eT

# Express as ~percent (dgdp is Δlog ⇒ approx growth rate), rounded to two decimal places
yT1.b = round(yT1.b * 100, 2)

# 1-step-ahead point forecast:
cat("ARIMA(1,1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is", as.numeric(yT1.b), "%.")
```

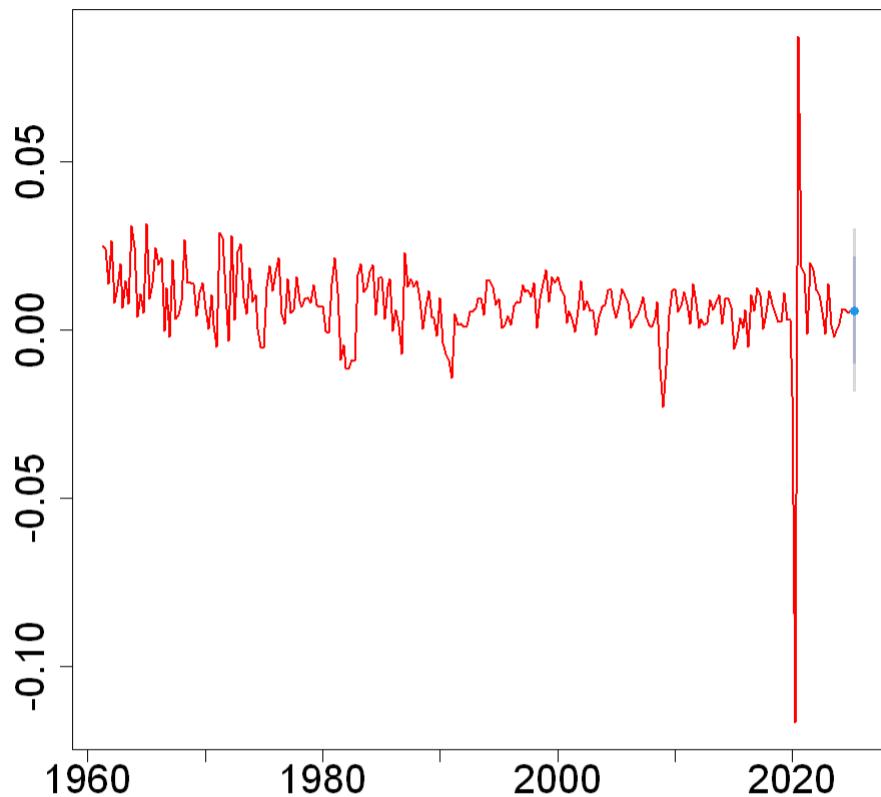
ARIMA(1,1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is 0.57 %.

In [40]:

```
# Set Jupyter/IRkernel figure size in inches for the output cell
options(repr.plot.width = 8, repr.plot.height = 8)

# Enlarge base R graphics text: axis labels, tick labels, and (potential) main title
par(cex.lab = 2, cex.axis = 2, cex.main = 2)

# Plot the 1-step-ahead forecast object from the 'forecast' package.
plot(f.b, aspect.ratio = 1, col = "red", lwd = 2, main = "")
```



## Recursive Out-of-Sample Forecasts for ARIMA Models

### What This Loop Is Doing (and Why)

We want to evaluate how an **AR(1)** forecasting model performs **out-of-sample**. To do this fairly, we mimic “real time”:

**1. Fix an initial estimation window** of size `Tstar`.

*Reason:* we need a reasonably long sample to estimate the AR(1) parameters reliably before we start judging forecasts.

**2. Expanding window (a.k.a. recursive estimation):**

At iteration `i`, we estimate the model using `dgdp[1 : (Tstar + i - 1)]`.

This mimics being at time `Tstar + i - 1` in real time, only using data **available then**.

**3. Forecast one step ahead ( $h = 1$ ):**

Using the model estimated at iteration `i`, we forecast the next observation, `dgdp[Tstar + i]`.

**4. Record results:**

We store the point forecast, its 95% interval, the realized value, and the absolute error.

**5. Repeat** until the end of the sample.

The number of iterations is `Tmax = T - Tstar`.

This produces a sequence of **genuine pseudo-real-time forecasts** for a fair accuracy assessment.

**Why not use the full sample each time?** Because that would leak future information and overstate forecast performance.

**Why not keep the window rolling with a fixed length?** That’s a different design (rolling/sliding window). Here we prefer the **expanding window** to use all past information accumulated up to each forecast origin, which often stabilizes AR parameter estimates.

## Recursive Out-of-Sample Forecasts for AR(1) model

```
In [22]: # https://www.rdocumentation.org/packages/forecast/versions/8.16/topics/forecast
library(forecast)

# Length of sample
T = length(dgdp)

# Container for results
results = data.frame()

# Initial estimation window
Tstar = 200

# Number of iterations
Tmax = T - Tstar

cat("Our data has", T, "observations.\n")
cat("We will keep the first", Tstar, "observations fixed for stability (initial estimation window).\n")
cat("From there, we add one observation at a time and forecast one step ahead at each iteration.\n")
cat("We will produce", Tmax, "out-of-sample 1-step-ahead forecasts, ending at observation", T, ".\n\n")

for (i in 1:Tmax) {
```

```

# Expand window estimation: use data up to time Tstar + i - 1
model = arima(dgdp[1:(Tstar + i - 1)], order = c(1, 0, 0))

# One-step-ahead forecast with a single 95% PI
f = forecast::forecast(model, h = 1, level = 95)

# Extract components (coerce to numeric to avoid ts attributes in arithmetic)
f.mean = as.numeric(f$mean[1]) # point forecast  $y_{T+1/T}$ 
f.lower95 = as.numeric(f$lower[1]) # 95% lower bound
f.upper95 = as.numeric(f$upper[1]) # 95% upper bound
obs.val = as.numeric(dgdp[Tstar + i]) # actual  $y_{T+1}$ 

# Absolute forecast error
f.error = abs(f.mean - obs.val)

# Row to append
temp = c(f.mean, f.lower95, f.upper95, obs.val, f.error)
results = rbind(results, temp)
}

# Name the columns
names(results) = c("Mean forecast", "Lower bound 95%", "Upper bound 95%", "Actual value", "Absolute error")

# Quick peek at the first 6 iterations
head(results)

# Quick peek at the last 6 iterations
tail(results)

# Save AR(1) rolling results
results.AR1 = results

# Mean Absolute Error (MAE)
mae.AR1 = mean(results.AR1[, "Absolute error"])

# Root Mean Squared Forecast Error (RMSFE)
rmsfe.AR1 = sqrt(mean((results.AR1[, "Mean forecast"] - results.AR1[, "Actual value"])^2))

cat("The MAE is", round(mae.AR1, 5), "and the RMSFE is", round(rmsfe.AR1, 5))

```

Our data has 256 observations.

We will keep the first 200 observations fixed for stability (initial estimation window). From there, we add one observation at a time and forecast one step ahead at each iteration. We will produce 56 out-of-sample 1-step-ahead forecasts, ending at observation 256 .

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.008011533	-0.007919723	0.02394279	0.0018136394	0.006197894
2	0.005725205	-0.010189328	0.02163974	0.0135896277	0.007864423
3	0.010346690	-0.005565246	0.02625863	0.0079159116	0.002430778
4	0.008136847	-0.007739376	0.02401307	0.0006311038	0.007505743
5	0.005283030	-0.010587519	0.02115358	0.0032454382	0.002037592
6	0.006279366	-0.009554827	0.02211356	0.0013675060	0.004911860

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
51	0.007309076	-0.01745116	0.03206931	-0.0003829651	0.007692042
52	0.007308181	-0.01742088	0.03203724	0.0017659246	0.005542256
53	0.007335084	-0.01735429	0.03202446	0.0062945465	0.001040538
54	0.007451791	-0.01718907	0.03209266	0.0059438355	0.001507956
55	0.007436371	-0.01715664	0.03202938	0.0050857420	0.002350629
56	0.007403866	-0.01714256	0.03195029	0.0054794362	0.001924430

The MAE is 0.00988 and the RMSFE is 0.02609

## Recursive out-of-sample forecasts for ARIMA(1,1) model

Note that the only difference between the two recursive loops (AR(1) and ARIMA(1,1)) is the order argument of the arima function.

```

In [23]: # https://www.rdocumentation.org/packages/forecast VERSIONS 8.16/topics/forecast
library(forecast)

# Length of sample
T = length(dgdp)

# Container for results
results = data.frame()

```

```

# Initial estimation window
Tstar = 200

# Number of iterations
Tmax = T - Tstar

cat("Our data has", T, "observations.\n")
cat("We will keep the first", Tstar, "observations fixed for stability (initial estimation window).\n")
cat("From there, we add one observation at a time and forecast one step ahead at each iteration.\n")
cat("We will produce", Tmax, "out-of-sample 1-step-ahead forecasts, ending at observation", T, ".\n\n")

for (i in 1:Tmax) {
  # Expand window estimation: use data up to time Tstar + i - 1
  model = arima(dgdp[1:(Tstar + i - 1)], order = c(1, 0, 1))

  # One-step-ahead forecast with a single 95% PI
  f = forecast::forecast(model, h = 1, level = 95)

  # Extract components (coerce to numeric to avoid ts attributes in arithmetic)
  f.mean = as.numeric(f$mean[1])                      # point forecast  $y_{T+1/T}$ 
  f.lower95 = as.numeric(f$lower[1])                    # 95% Lower bound
  f.upper95 = as.numeric(f$upper[1])                   # 95% upper bound
  obs.val = as.numeric(dgdp[Tstar + i])                # actual  $y_{T+1}$ 

  # Absolute forecast error
  f.error = abs(f.mean - obs.val)

  # Row to append
  temp = c(f.mean, f.lower95, f.upper95, obs.val, f.error)
  results = rbind(results, temp)
}

# Name the columns
names(results) = c("Mean forecast", "Lower bound 95%", "Upper bound 95%", "Actual value", "Absolute error")

# Quick peek at the first 6 interations
head(results)

# Quick peek at the last 6 iterations
tail(results)

# Save ARMA(1,1) rolling results
results.ARMA11 = results

# Mean Absolute Error (MAE)
mae.ARMA11 = mean(results.ARMA11[, "Absolute error"])

# Root Mean Squared Forecast Error (RMSFE)
rmsfe.ARMA11 = sqrt(mean((results.ARMA11[, "Mean forecast"] - results.ARMA11[, "Actual value"])^2))

cat("The MAE is", round(mae.ARMA11, 5), "and the RMSFE is", round(rmsfe.ARMA11, 5))

```

Our data has 256 observations.

We will keep the first 200 observations fixed for stability (initial estimation window).  
 From there, we add one observation at a time and forecast one step ahead at each iteration.  
 We will produce 56 out-of-sample 1-step-ahead forecasts, ending at observation 256 .

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.008304613	-0.007545557	0.02415478	0.0018136394	0.006490974
2	0.006067570	-0.009768451	0.02190359	0.0135896277	0.007522058
3	0.009450258	-0.006380321	0.02528084	0.0079159116	0.001534347
4	0.008524317	-0.007268620	0.02431725	0.0006311038	0.007893213
5	0.005759379	-0.010031834	0.02155059	0.0032454382	0.002513941
6	0.005805135	-0.009951210	0.02156148	0.0013675060	0.004437629

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
51	0.005975644	-0.01855346	0.03050475	-0.0003829651	0.0063586096
52	0.005740100	-0.01874974	0.03022994	0.0017659246	0.0039741757
53	0.005631166	-0.01881470	0.03007703	0.0062945465	0.0006633802
54	0.005654663	-0.01874172	0.03005105	0.0059438355	0.0002891724
55	0.005706781	-0.01864266	0.03005622	0.0050857420	0.0006210393
56	0.005694128	-0.01860657	0.02999483	0.0054794362	0.0002146917

The MAE is 0.00984 and the RMSFE is 0.02605

## Conclusion

Both AR(1) and ARIMA(1,1) models perform virtually identical when comes to forecasting growth rate of the Canadian GDP one period ahead.

## Forecasting ADL Models

### ADL(0,1) Model

**Note:** We will forecast **ADL(0,1)** and **ADL(1,1)** models for illustration, so we can match R's built-in functions with our manual computations. As we saw last week, the **best model by BIC** was **ADL(0,1)**, while the **best model by AIC** was a **subset ADL** with the first and third lag of the change in weekly hours worked.

**R:** Forecasting GDP growth rate using the lagged value of the change in hours worked, 1-step ahead:

$$y_t = \alpha_0 + \beta_1 x_{t-1} + \varepsilon_t$$

```
In [24]: # Align series on overlapping dates (common time index)
```

```
data = ts.intersect(dgdp, dhours)
```

```
# ADL(0,1): dgdp_t on Lagged dhours_{t-1}
```

```
library(dynlm)
```

```
model.c = dynlm(dgdp ~ L(dhours, 1), data)
```

```
# Print a compact regression table
```

```
library(stargazer)
```

```
stargazer(model.c, type = "text")
```

```
=====
```

```
Dependent variable:
```

```
-----
```

```
dgdp
```

```
-----
```

```
L(dhours, 1) 0.005**
```

```
(0.002)
```

```
Constant 0.008***
```

```
(0.001)
```

```
-----
```

```
Observations 249
```

```
R2 0.022
```

```
Adjusted R2 0.018
```

```
Residual Std. Error 0.013 (df = 247)
```

```
F Statistic 5.617** (df = 1; 247)
```

```
=====
```

```
Note: *p<0.1; **p<0.05; ***p<0.01
```

```
In [25]:
```

```
# 1-step ahead forecast
```

```
newdata = data.frame("dhours" = tail(data[, "dhours"], n=1))
```

```
# 95% prediction interval (future outcome uncertainty)
```

```
f.c = predict(model.c, newdata = newdata, interval = "prediction", level = 0.95)
```

```
f.c
```

```
A matrix: 1 x 3 of type dbl
```

	fit	lwr	upr
1	0.008041573	-0.01667723	0.03276038

```
In [26]:
```

```
# --- Manual calculations for 1-step-ahead ADL(0,1) forecast from model.c ---
```

```
# https://www.econometrics-with-r.org/14.5-apatadlm.html
```

```
coef.c = coef(model.c)
```

```
alpha0 = coef.c["(Intercept)"]
```

```
beta1 = coef.c["L(dhours, 1)"]
```

```
# Use the most recent dhours (time T) to predict dgdp at T+1
```

```
T = nrow(data)
```

```
yT1.c = alpha0 + beta1 * data[T, "dhours"]
```

```
# Express as ~percent (dgdp is ΔLog ⇒ approx growth rate), rounded to two decimal places
```

```
yT1.c = round(yT1.c * 100, 7)
```

```
# 1-step-ahead point forecast:
```

```
cat("ADL(0,1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is", as.numeric(yT1.c), "%.")
```

```
ADL(0,1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is 0.8041573 %.
```

## Forecasting ADL Models

## ADL(1,1) Model

R: Forecasting GDP growth rate using the **ADL(1,1)** model, 1-step ahead:

$$y_t = \alpha_0 + \alpha_1 y_{t-1} + \beta_1 x_{t-1} + \varepsilon_t$$

```
In [27]: # Align series on overlapping dates (common time index)
data = ts.intersect(dgdp, dhours)
```

```
# ADL(1,1): dgdp_t on Lagged dgdp_{t-1} and Lagged dhours_{t-1}
library(dynlm)
model.d = dynlm(dgdp ~ L(dgdp, 1) + L(dhours, 1), data)
```

```
# Print a compact regression table
library(stargazer)
stargazer(model.d, type = "text")
```

```
=====
Dependent variable:
-----
dgdp
-----
L(dgdp, 1)           -0.016
                   (0.065)

L(dhours, 1)          0.005**
                   (0.002)

Constant              0.008***
                   (0.001)

-----
Observations          249
R2                    0.022
Adjusted R2           0.015
Residual Std. Error   0.013 (df = 246)
F Statistic            2.828* (df = 2; 246)
-----
Note:                 *p<0.1; **p<0.05; ***p<0.01
```

```
In [28]: # 1-step ahead forecast
newdata = data.frame("dgdp" = tail(data[, "dgdp"], n=1), "dhours" = tail(data[, "dhours"], n=1))
# 95% prediction interval (future outcome uncertainty)
f.d = predict(model.d, newdata = newdata, interval = "prediction", level = 0.95)
f.d
```

```
A matrix: 1 × 3 of type dbl
      fit      lwr      upr
1 0.008216719 -0.01658865 0.03302209
```

```
In [29]: # --- Manual calculations for 1-step-ahead ADL(0,1) forecast from model.c ---
# https://www.econometrics-with-r.org/14.5-apatadlm.html
coef.d = coef(model.d)

alpha0 = coef.d["(Intercept)"]
beta1 = coef.d["L(dhours, 1)"]
alpha1 = coef.d["L(dgdp, 1)"]

#Use the most recent dgdp and the most recent dhours (time T) to predict dgdp at T+1
T = nrow(data)
yT1.d = alpha0 + alpha1 * data[T, "dgdp"] + beta1 * data[T, "dhours"]

# Express as ~percent (dgdp is Δlog ⇒ approx growth rate), rounded to two decimal places
yT1.d = round(yT1.d * 100, 7)

# 1-step-ahead point forecast:
cat("ADL(1,1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is", as.numeric(yT1.d), "%.")
```

```
ADL(1,1) generated 1-step-ahead forecast of the growth rate of the Canadian GDP is 0.8216719 %.
```

## Recursive Out-of-Sample Forecasts for ADL(0,1) Model

```
In [30]: # Align series on overlapping dates (common time index)
data = ts.intersect(dgdp, dhours)
```

```
# Length of sample
T = nrow(data)
```

```
# Container for results
results = data.frame()
```

```
# Initial estimation window
Tstar = 200
```

```
# Number of iterations
```

```

Tmax = T - Tstar

for (i in 1: (T - Tstar)) {
  model = dynlm(dgdp ~ L(dhours, 1), data[1: (Tstar + i - 1), ])
  newdata = data.frame("dhours" = data[(Tstar + i - 1), "dhours"])

  # One-step-ahead forecast with a single 95% PI
  f = predict(model, newdata = newdata, interval = "prediction", level = 0.95)

  # Extract components (coerce to numeric to avoid ts attributes in arithmetic)
  f.mean = as.numeric(f[1,1])
  f.lower95 = as.numeric(f[1,2])
  f.upper95 = as.numeric(f[1,3])

  # Realized value at time (Tstar + i) to compare against the forecast
  obs.val = as.numeric(data[(Tstar + i), "dgdp"])

  # Absolute forecast error
  f.error = abs(f.mean - obs.val)

  # Row to append
  temp = c(f.mean, f.lower95, f.upper95, obs.val, f.error)
  results = rbind(results, temp)
}

names(results)=c("Mean forecast", "Lower bound 95%", "Upper bound 95%", "Actual value", "Absolute error")

# Quick peek at the first 6 iterations
head(results)

# Quick peek at the last 6 iterations
tail(results)

# Save ADL(0,1) rolling results
results.ADL01 = results

# Mean Absolute Error (MAE)
mae.ADL01 = mean(results.ADL01[, "Absolute error"])

# Root Mean Squared Forecast Error (RMSFE)
rmsfe.ADL01 = sqrt(mean((results.ADL01[, "Mean forecast"] - results.ADL01[, "Actual value"])^2))

cat("The MAE is", round(mae.ADL01, 5), "and the RMSFE is", round(rmsfe.ADL01, 5))

```

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.011759182	-0.01319617	0.03671454	0.0018136394	0.009945543
2	0.002835183	-0.02217004	0.02784041	0.0135896277	0.010754445
3	0.009134477	-0.01561688	0.03388583	0.0079159116	0.001218565
4	0.010184359	-0.01462690	0.03499561	0.0006311038	0.009553255
5	0.008609535	-0.01612281	0.03334188	0.0032454382	0.005364097
6	0.005984830	-0.01876233	0.03073199	0.0013675060	0.004617324

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
45	0.008609535	-0.01612281	0.03334188	0.010229149	0.001619613
46	0.007034712	-0.01768454	0.03175396	0.005138201	0.001896511
47	0.008609535	-0.01612281	0.03334188	-0.001370987	0.009980522
48	0.007034712	-0.01768454	0.03175396	0.013454326	0.006419614
49	0.009659418	-0.01511825	0.03443709	0.001425423	0.008233995
50	0.008609535	-0.01612281	0.03334188	-0.002258095	0.010867631

The MAE is 0.00962 and the RMSFE is 0.02192

## Recursive Out-of-Sample Dorecasts for ADL(1,1) Model

```

In [31]: # Align series on overlapping dates (common time index)
data = ts.intersect(dgdp, dhours)

# Length of sample
T = nrow(data)

# Container for results
results = data.frame()

# Initial estimation window
Tstar = 200

```

```

# Number of iterations
Tmax = T - Tstar

for (i in 1: (T - Tstar)) {
  model = dynlm(dgdp ~ L(dgdp, 1) + L(dhours, 1), data[1: (Tstar + i - 1), ])
  newdata = data.frame("dgdp" = data[(Tstar + i - 1), "dgdp"],
                       "dhours" = data[(Tstar + i - 1), "dhours"])

  # One-step-ahead forecast with a single 95% PI
  f = predict(model, newdata = newdata, interval = "prediction", level = 0.95)

  # Extract components (coerce to numeric to avoid ts attributes in arithmetic)
  f.mean = as.numeric(f[1,1])
  f.lower95 = as.numeric(f[1,2])
  f.upper95 = as.numeric(f[1,3])

  # Realized value at time (Tstar + i) to compare against the forecast
  obs.val = as.numeric(data[(Tstar + i), "dgdp"])

  # Absolute forecast error
  f.error = abs(f.mean - obs.val)

  # Row to append
  temp = c(f.mean, f.lower95, f.upper95, obs.val, f.error)
  results = rbind(results, temp)
}

names(results)=c("Mean forecast", "Lower bound 95%", "Upper bound 95%", "Actual value", "Absolute error")

# Quick peek at the first 6 iterations
head(results)

# Quick peek at the last 6 iterations
tail(results)

# Save ADL(1,1) rolling results
results.ADL11 = results

# Mean Absolute Error (MAE)
mae.ADL11 = mean(results.ADL11[, "Absolute error"])

# Root Mean Squared Forecast Error (RMSFE)
rmsfe.ADL11 = sqrt(mean((results.ADL11[, "Mean forecast"] - results.ADL11[, "Actual value"])^2))

cat("The MAE is", round(mae.ADL11, 5), "and the RMSFE is", round(rmsfe.ADL11, 5))

```

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
1	0.011674797	-0.01331729	0.03666688	0.0018136394	0.009861158
2	0.002872213	-0.02215299	0.02789742	0.0135896277	0.010717414
3	0.008980548	-0.01579286	0.03375396	0.0079159116	0.001064637
4	0.010102980	-0.01473300	0.03493896	0.0006311038	0.009471876
5	0.008635228	-0.01614105	0.03341151	0.0032454382	0.005389789
6	0.005987289	-0.01877879	0.03075337	0.0013675060	0.004619783

A data.frame: 6 × 5

	Mean forecast	Lower bound 95%	Upper bound 95%	Actual value	Absolute error
	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
45	0.008482817	-0.01626973	0.03323536	0.010229149	0.001746331
46	0.006936716	-0.01780460	0.03167803	0.005138201	0.001798516
47	0.008573624	-0.01618317	0.03333042	-0.001370987	0.009944610
48	0.007095270	-0.01766379	0.03185433	0.013454326	0.006359057
49	0.009504839	-0.01529319	0.03430287	0.001425423	0.008079416
50	0.008624371	-0.01614750	0.03339624	-0.002258095	0.010882466

The MAE is 0.00961 and the RMSFE is 0.02184

## Comparison Across the Models

AR(1), ARMA(1,1), ADL(0,1), ADL(1,1)

In [32]:

```

data.frame("AR(1)" = mae.AR1, "ARMA(1,1)" = mae.ARMA11, "ADL(0,1)" = mae.ADLO1, "ADL(1,1)" = mae.ADL11)
data.frame("AR(1)" = rmsfe.AR1, "ARMA(1,1)" = rmsfe.ARMA11, "ADL(0,1)" = rmsfe.ADLO1, "ADL(1,1)" = rmsfe.ADL11)

```

A data.frame: 1 × 4

**AR.1. ARMA.1.1. ADL.0.1. ADL.1.1.**

<dbl>	<dbl>	<dbl>	<dbl>
-------	-------	-------	-------

0.009880319	0.009840262	0.009619095	0.00961235
-------------	-------------	-------------	------------

A data.frame: 1 × 4

**AR.1. ARMA.1.1. ADL.0.1. ADL.1.1.**

<dbl>	<dbl>	<dbl>	<dbl>
-------	-------	-------	-------

0.02608751	0.02605203	0.02192396	0.02184136
------------	------------	------------	------------

Our tables shows the lowest MAE and RMSFE for ADL(1,1), slightly better than ADL(0,1) and clearly better than AR(1)/ARMA(1,1). That means ADL(1,1) is your best of the four on this sample.

In [ ]: